

---

# Hijacking Simulators with Universal Probabilistic Programming

---

Paper ID: 1010

## Abstract

Notes on how to develop a formalism to perform inference in population based simulators.

## 1 Introduction

One challenge in performing inference in such models is the combinatorial increase in *paths* that can be taken within the given probabilistic model. This is because each member of the population has a bounded, but large number of decisions to make. In addition to this, each member of the population can interact with other members of the population. To model a single population member quantum-based methods leveraging linear superposition and entanglement provide an ability to explore multiple paths simultaneously maintaining the whole program, *state*, in one iteration. Unfortunately, analytically writing the functional form of the state is, in this instance, impossible. The usability of current classical inference methods in this context is also problematic. However, by combining hijacking and amortizing parts of the simulator, in particular rejection sampling loops, we can not only correctly perform inference in such models, as the rejection sampling loops are now replaced with learnt functions, but we can also perform inference much more efficiently. This is because the rejection sampling loops are now replaced with learnt functions which enables us to correctly construct the density and as they are amortized, evaluation is cheap and only has to be performed once per call. This significantly reduces memory consumption and running time.

Without the amortization procedure, as updates are made sequentially MCMC methods are typically not-well suited this type of recursive estimation problem. As each time we get a new data point  $y_{t+1}$  we have to run new MCMC simulations for  $P(x_{0:t_1} | y_{0:t+1})$ , we also cannot directly reuse the samples  $\{x_{0:t}^i\}$ . However, as we can only run the simulator `forward()` we cannot return back to a previous state and re-run. So adapting existing inference techniques is the only viable option. MCMC methods can be effective with a well designed amortized inference procedure (**Note: This is not guaranteed**). Non-MCMC based methods such as importance sampling (IS) and sequential monte carlo (SMC), may also be sensible inference engines, although we shall discuss the limitations within the simulator context in the proceeding sections.

However, to implement this in practice we are required to do the following:

- Create a sensible training procedure
- Have 2 additional functions implemented within `pyprob`
  1. 1) That stores only the necessary raw samples from the desired rejection sampling addresses to be used for training.
  2. 2) A function at *test time* that replaces those given addresses with a learnt surrogate.
- Finer grained control over the actual latent variables of interest, to avoid unnecessary memory consumption and ensure that inference problems remain tractable.

There is no doubt that implementing some of these features will be more invasive to the underlying simulator code than previously just hijacking the raw-random number draws, but, it would enable

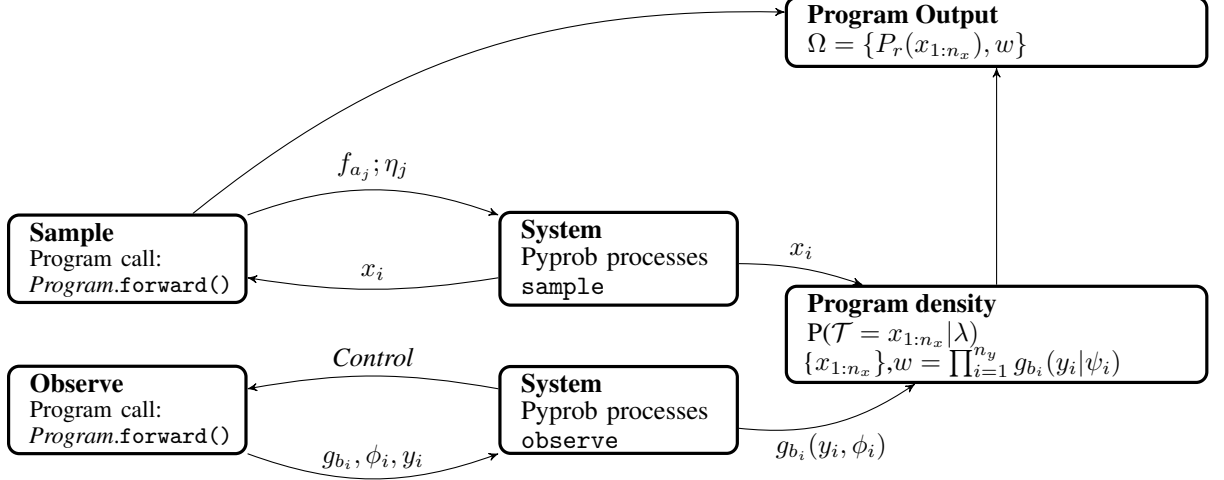


Figure 1: An overview of the hijacking system.  $x_i$  represent latent variables.  $f_{a_i}$  represent raw sample calls,  $g_{b_j}$  represent conditioning statements.  $\mathcal{T}$  corresponds to the program trace.  $w$  represents the collective weights.  $r$  represents run number.  $n_x$  represents the number of latent variables and  $n_y$  represents the number of observations, which will always be fixed, but is model dependent.  $\lambda$  represents all the other variables generated by running the simulator `.forward()`.

us to attack difficulties in the inference procedures that are currently not amenable to the current approach.

## 2 Limitations with different inference engines

### 2.1 Prior based sampling

In prior based sampling we take our existing program and look directly at the product of all the raw sample calls, denoted  $f_{a_i}(x_i|\eta)$  and construct a proposal that is dependent on those. In this sampling scheme our proposal is defined as:

$$q(x_{1:n_x}) = \begin{cases} \prod_{i=1}^{n_x} f_{a_i}(x_i|\eta) & \text{if } \mathcal{B}(x_{1:n_x}) = 1 \\ 0 & \text{otherwise} \end{cases} \quad (1)$$

where  $\mathcal{B}(\cdot)$  is satisfied by the simulator and  $f_{a_i}(\cdot)$  represents a raw random sample,  $\eta$  is a subset of the variables in-scope at the point of sampling and  $n_x$  is the number of latent variables. This is true in all inference engines used in this hijacking framework. However, this makes use of no conditioning and is not sufficient for performing inference in complex, population-based simulations.

## 3 Non-prior Based Sampling

In non-prior based sampling we utilise conditioning too, which enables us to explore spaces more efficiently. In the current hijacking set-up, we can only do conditioning at the end of each `forward()` run of the simulator. This does inhibit certain inference engines and makes aspects of performing inference within this setting more complicated. Nonetheless, there are several inference schemes that we can exploit and utilise. We shall now give an Introduction to each of these schemes.

### 3.1 Importance Sampling

In the generalised probabilistic programming setting an importance sampling scheme over the program raw samples and observations can be defined as follows:

$$p(\mathcal{T} = x_{1:n_x} | \lambda) = \begin{cases} \prod_{i=1}^{n_x} f_{a_i}(x_i|\eta) \prod_{j=1}^{n_y} g_{b_j}(y_j|\phi) & \text{if } \mathcal{B}(x_{1:n_x}) = 1 \\ 0 & \text{otherwise} \end{cases} \quad (2)$$

where the observe statements are defined by  $g_{b_j}(\cdot)$ ,  $\phi$  plays the same role as  $\eta$  and  $n_y$  represents the number of observations. The expectation over the whole program can be defined as:

$$\mathbb{E}_{p(\Omega|\lambda)}[h(\Omega)] = \mathbb{E}_{p(x_{1:n_x}|\lambda)}[h(P_r(x_{1:n_x}))] \quad (3)$$

**Note: It is a little weird how we define this, as the the program output  $\Omega$  should also include the product of the weights**

where  $r$  represents a run of the program,  $P_r(\cdot)$  is the entirety of the program output,  $h(\cdot)$  is just a deterministic function.

As the importance sampler updates the weights according to the following formula:

$$w(x_{1:n_x}) = \prod_{i=1}^{n_x} \frac{f_{a_i}(x_i|\eta_i)}{q_i(x_i|f_{a_i}, \eta_i)} \prod_{j=1}^{n_y} g_{b_j}(y_j|\phi_j) \quad (4)$$

When the number of observations is fixed, the combined weight is equal to  $w(x_{1:n_x}) = \prod_{j=1}^{n_y} w_j$ . To generate the weight up to each observe we can define  $K(j)$  to be a function that returns the least  $i$  before hitting an observe  $j$ . This enables us to write the  $j$ -th weight as follows:

$$w_j(x_1 : K(j)) = g_{b_j}(y_j|\psi_j) \prod_{K(j-1)+1}^{K(j)} v_i \quad (5)$$

where we define  $v_j = \frac{f_{a_j}(x_j|\eta_j)}{q_j(x_j|f_{a_j}, \eta_j)}$

Taking these definitions we can construct a target density over the latents that is proportional to the products of sample and observe statements **Note: Add update posterior here**. Thus, given the proposal  $q(x_{1:K(t)})$  we define the importance weight as:

$$w(x_{1:t}) = \prod_{j=1}^t w_j(x_{1:K(j)}) = \frac{p_t(x_{1:K(t)})}{q(x_{1:K(t)})} = \begin{cases} \prod_{j=1}^t g_{b_j}(y_j|w_i) \prod_{K(j-1)+1}^{K(j)} v_i & \text{if } B_t(x_1 : K(t)) = 1 \\ 0 & \end{cases} \quad (6)$$

One problem with importance sampling is in its ability to scale and condition upon multiple observations. | Choosing a *good* prior remains challenging.

### 3.2 Sequential Monte Carlo

Given multiple observations, we can ake the inference more computationally efficient by employing sequential monte carlo (SMC), although as we can only condition at the end of one full run of the simulation, computational speed increases are minimal. SMC allows for conditioning on multiple observations, but the number of observations must be fixed. **Note: If we could condition on observations within the simulator, we could get dramatic speed increases over standard importance sampling.**

To implement SMC in the current set-up we round require a forking method and the

### 3.3 Random-walk Metropolis Hastings

Due to the current set-up performing Random-walk Metropolis Hastings (RMH) is not feasible as the rejection sampling blocks within the simulator mean that we cannot correctly perform RMH as we do not have a functional form for all parts of the simulator to construct a target density that we can evaluate. **Note: Ask tom what: one can sometimes achieve improvements by instead using the type of the distribution object associated with  $x_m$  (what does this first part mean?) to automatically construct a valid random walk proposal that allows for improved hill climbing behavior on the individual updates. .**

---

**Algorithm 1** SMC for Population Based Simulators - well any general program - not new
 

---

```

1: procedure SMC( $T, M, \mathbf{x}, \mathbf{w}, \mathbf{v}, \mathbf{g}$ )
2:   for  $t = 1 : T$  do
3:     for  $m = 1 : M$  do
4:        $w_t^m = 1$ 
5:       for  $k = K(t-1) + 1 : K(t)$  do
6:          $x_k^m \sim q_k(x_k^m | f_{a_k^m}, \eta_j^m)$ 
7:          $w_t^m \leftarrow w_t^m v_j^m(x_k^m)$ 
8:       end for
9:     end for
10:     $w_t^m \leftarrow w_t^m g_{b_t^m}(y_t^m | \psi_t^m)$ 
11:     $\{x_{1:t}^{1:m}\} \leftarrow \text{resample}(\{x_{1:t}^{1:m}\}, \{w_{1:t}^{1:m}\})$ 
12:  end for
13: end procedure

```

---




Figure 2: The DAG of the Parasite model.

## 4 An Example: Modelling Parasite Densities in Simulation

The parasite model by Smith et al. [5] is a simple stochastic model for determining the number of infections within a given population demographic. The models hypothesized in [5] were found, after extensive testing and development, to not characterise well the number of infections among different demographics, especially those under 15 years old. By using probabilistic programming we can immediately understand the expressiveness of the proposed model and can extract important insights for extending the proposed models, which increases ones ability to develop improved models to estimate parasite densities, given our observations of the seasonal entomological inoculation rate (EIR). In addition to this we can greatly reduce the computational costs, whilst improving model comprehension and expressivity.

**Common features of the models** The models used are all discrete time micro-simulations of malaria in humans, based on the published model used in [1], which we refer to as the base model. The model for infection of the vector as a function of parasite densities [2] (and hence for the effect of vaccination on onward transmission), the case-management model [3], the model of vaccination by pre-erythrocytic vaccines [4], the models for clinical outcomes and mortality [5,6], and the basic design of the scenarios used for predicting impacts [3,7] all correspond to the previous descriptions, which are accompanied by detailed rationales for the formulations used.

For all models, the number of infections introduced in individual  $i$  in five-day time step  $t$ , is distributed as:

$$h(i, t) \sim \text{Pos}(\lambda(i, t)) \quad (7)$$

where the expected number  $\lambda(i, t) = S(i, t)E_a(i, t)$  and  $S(i, t)$  is the susceptibility (proportion of inoculations that result in infections) and  $E_a(i, t)$  is the expected number of entomological inoculations, adjusted for age and individual factors.

**Base model** Although a statistical model is not formally defined within [5] we define the unknown *disease dynamics* parameters with a set of latent variables  $\{x_i\}^N$  and construct the model as follows:

In the base model  $E_a(i, t) = E_{max}(t) \frac{A(a(i, t))}{A_{max}}$  where  $E_{max}(t)$  refers to the usual measure of the EIR computed from human bait collections on adults  $A(\cdot)$  and, the availability to mosquitoes, is assumed to be proportional to average body surface area and to depend only on age  $a(i, t)$ . The function used to compute  $A(a(i, t))$  increases with age up to age 20 years where it reaches a value of  $A_{max}$ .

To capture observed relationships between infection and measured entomological exposure, the susceptibility is then:

$$S(i, t) = (S_\infty + \frac{1 - S_\infty}{1 + \frac{E_a(i, t)}{E^*}})(S_{imm} + \frac{1 - S_{imm}}{1 + (\frac{X_p(i, t)}{X_p^*})^{\gamma_p}}) \quad (8)$$

where  $S_{imm}, X_p^*, E^*, \gamma_p, S_\infty$  are latent variables and  $X_p(i, t) = \int_{-a(i, t)}^t E_a(i, \tau) d\tau$ . Fitting this model to data of [2] using their method learns the values  $S_\infty = 0.049$  and  $E^* = 0.032$  inoculations/person-night. The model for each individual infection  $j$  in host  $i$  comprises a time series of parasite densities. In a naïve host the expectation of the logarithm of the density of the infection at time point  $\tau$  in the course of the infection,  $\ln(y_0(i, j, \tau))$ , is taken from a statistical description of parasite densities in malaria therapy patients [3]. In the presence of previous exposure and co-infection, the expected log density for each concurrent infection is:

$$E(\ln(y(i, j, \tau))) = D_y(i, t)D_m(i, t)D_h(i, t)\ln(y_0(i, j, \tau)) + \ln(\frac{D_x}{M(i, t)} + 1 - D_x) \quad (9)$$

where  $M(i, t)$  is the total multiplicity of infection and

$$D_y(i, t) = \frac{1}{1 + \frac{X_y(i, j, t)}{X_y^*}} \quad (10)$$

where  $X_y(i, j, t) = \sum_{t-a}^t Y(i, t) - \sum_{t_0, j}^t y(i, j, \tau)$  (note that a continuous time approximation to this is given in the original publications [4, 3] and hence measures the cumulative parasite load, and **TODO: To finish**

**TODO: Add some of the plots generated from this model and an accompanying analysis to them.**

## 5 Amortized Rejection Sampling

One common problem in traditional simulators, *programs*, outside of the standard probabilistic programming framework is that they rely heavily on rejection sampling to generate samples from the target density. Not only is rejection sampling computationally inefficient, [especially if criteria in which samples are generated has been poorly calibrated] (may leave out this statement as it may make a reviewer question why would amortization work any better.), but it creates complications for inference. As the rejection sampling loops provide no explicit form of the target density, we can not simply evaluate the density and leverage statistically and computationally efficient algorithms such as MCMC sampling. This inhibits our ability to transform an arbitrary stochastic simulator into a probabilistic programming framework [1]. To this end, we derive both a mathematical and engineering based framework to facilitate the amortization of rejection sampling loops in arbitrary simulators, which provides a path way to both leverage a larger, more efficient class of inference methods and convert arbitrary stochastic simulators into probabilistic programs systems.

## 6 Details

We do this by Replacing the rejection sampling loops with learnt functions from an amortized inference procedure Figure 3. This then enables us to replace rejection sampling loops with amortized functions that are quick to evaluate, provide a functional form to the density and significantly reduce memory over-heads, as compared to [1], due to not having to track the thousands of raw sample calls  $f_{a_i}$  generated in standard simulators.

Consider the following program, that is representative of a simple stochastic simulator and two types of rejection sampler:




Figure 3: Using PyProb we can locate all regions of inefficient sampling and replace those points in the code with learnt functions that represent the distributions of the rejection sampling statements.  $Z_i$  represents the marginal and by learning surrogate functions for each of the rejection sampling loops. In doing so we can learn approximate raw sample distributions  $f(x_i|\eta)$  for each rejection sampling loop. This means that we can evaluate the function at each part in the simulator, enabling us to construct a target density that can be evaluated when using MCMC inference schemes such as RMH.

Stochastic Simulator	Rejection Sampler 1	Rejection Sampler 2
<pre> def f():     z1 ~ N(0, 1)     z2 ← R1(z1, 1)     z3 ~ U(0, 2)     z4 ← R2(z2, z3, 1)      return g(z1:4) </pre>	<pre> def R1(z1):     x ~ N(z1, 1)     if x &gt; 0:         return x     else:         return R1(z1) </pre>	<pre> def R2(z2, z3):     y ← 0     while y &lt; z2:         y ~ N(z3, 1)     end     return y </pre>
$P(\mathcal{T} = x_{1:n_x}   \lambda)$	$q_1(z_2   z_1)$	$q_2(z_4   z_2, z_3)$

Here the rejection sampling functions take an additional argument 1, or 0, to specify whether or not we want to collect samples to train, 1, else replace with the learnt function, 0.  $P$  represents the program density and  $q_1$  and  $q_2$  represents the unknown densities that we wish to learn through an amortized approach.

However, in order to do this, we must correctly specify our variational objective.

## 7 Constructing the Objective

The framework we present in this section is generally applicable to stochastic simulator environments, where there exists rejection sampling loops.

We start by noting that the prior, defined in the context of the outer samples is defined as:

$$p(x) = \begin{cases} \prod_{j=1}^{n_x} f_{a_j}(x_j; \psi_j) & \text{if } \mathcal{B}(x) = 1 \\ 0 & \text{else} \end{cases} \quad (11)$$

We start with the information projection we construct the variational objective as follows:

$$\begin{aligned}
J(\eta) &= KL(p(x) || q(x; \eta)) \\
&= \int p(x) \ln \left( \frac{p(x)}{q(x; \eta)} \right) dx \\
&\quad \text{Let } a_j \in \{1, \dots, L\}, S_R \subset \{1, \dots, L\} \\
&= \mathbb{E}_{p(x)} \left[ \sum_{j=1, a_j \notin S_R}^{n_x} \ln(p(x)) \right] - \mathbb{E}_{p(x)} \left[ \sum_{j=1, a_j \in S_R}^{n_x} \ln(q_{a_j}(x_j; \eta)) \right]
\end{aligned}$$

no dependence on  $\eta$

Thus, to define the objective for each subproblem  $r \in S_R$  as:

$$J_r(\eta_r) = -\mathbb{E}_{p(x)} \left[ \sum_{j=1}^{n_x} \mathbb{I}(a_j = r) \ln(q_r(x_j; \eta_r, \phi_j)) \right] \quad (12)$$

where during training we extract samples from each forward run and train each rejection sampler separately.

## 8 Additional TODOs

- Before implementing such tools into the simulator, spend time looking for simpler stochastic simulators that can be hijacked but have target densities that can be constructed analytically.
- If we are not utilizing the additional information stored in the trace, which is only used during the inference compilation training procedure, then we do not need it all, as for importance sampling we only require the weights after each `forward()` update.
- If we cannot find any pre-existing simulators that are amenable to this approach then, we need to create our own, with a rejection sampling loop, that can be exploited by the new machinery. This would enable us to debug, provide a simple example for the paper and understand the feasibility of the entire approach.

## References

- [1] Atilim Gunes Baydin, Lukas Heinrich, Wahid Bhimji, Bradley Gram-Hansen, Gilles Louppe, Lei Shao, Kyle Cranmer, Frank Wood, et al. Efficient probabilistic inference in the quest for physics beyond the standard model. *arXiv preprint arXiv:1807.07706*, 2018.
- [2] John C Beier, Charles N Oster, Fred K Onyango, James D Bales, James A Sherwood, Peter V Perkins, David K Chumo, Davy V Koech, Richard E Whitmire, Clifford R Roberts, et al. Plasmodium falciparum incidence relative to entomologic inoculation rates at a site proposed for testing malaria vaccines in western kenya. *The American journal of tropical medicine and hygiene*, 50(5):529–536, 1994.
- [3] Nicolas Maire, Thomas Smith, Amanda Ross, Seth Owusu-Agyei, Klaus Dietz, and Louis Molineaux. A model for natural immunity to asexual blood stages of plasmodium falciparum malaria in endemic areas. *The American journal of tropical medicine and hygiene*, 75(2\_suppl):19–31, 2006.
- [4] Thomas Smith, Gerry F Killeen, Nicolas Maire, Amanda Ross, Louis Molineaux, Fabrizio Tediosi, Guy Hutton, Jürg Utzinger, Klaus Dietz, and Marcel Tanner. Mathematical modeling of the impact of malaria vaccines on the clinical epidemiology and natural history of plasmodium falciparum malaria: Overview. *The American journal of tropical medicine and hygiene*, 75(2\_suppl):1–10, 2006.
- [5] Thomas Smith, Nicolas Maire, Klaus Dietz, Gerry F Killeen, Penelope Vounatsou, Louis Molineaux, and Marcel Tanner. Relationship between the entomologic inoculation rate and the force of infection for plasmodium falciparum malaria. *The American journal of tropical medicine and hygiene*, 75(2\_suppl):11–18, 2006.