
Hijacking Simulators with Universal Probabilistic Programming

Paper ID:

Abstract

Probabilistic programming provides a way to perform statistical inference over simulations of events in a programmatic way. Thus, simulators by design are ideal programs for probabilistic programming systems (PPSs). However, within existing PPSs one would have to re-implement the simulator via the PPS language specification, which is inefficient and often not feasible due to the complexity of such scientific and industrial simulators. However, recent work by [2] demonstrated a pathway to turn a particular type of event-based simulator into a probabilistic program. But, this still meant that a large class of critically important population-based simulators could not be turned into probabilistic programs and as such could not be used in a probabilistic programming framework. In this work, we extend that framework to encompass population-based simulators, a very large class of simulators that are used extensively across epidemiology modeling, multi-agent modeling and financial modeling. We enable simulators written in 13 different languages to be easily converted to probabilistic programs, for which we algorithmically outline how to do and provide examples of. We demonstrate our method on three models and provide an infected population simulator and demonstrate the hijacking on two important malaria population-based simulators OpenMalaria and EMOD.

1 Introduction

[BG: The audience is a probabilistic programming audience - so that is what we will target]

Probabilistic programming [13, 19] is a machine learning paradigm that enables users to write, and perform inference in, complex probabilistic models in a programmatic, yet efficient and simplified way. The source code of the given model, or equivalently the *program*, is then transformed into a program density, which can be viewed as the joint density of the model. Having direct access to the source code enables probabilistic programming system (PPS) to conduct several processes in an automated fashion, such as automating inference [29, 12, 28]. In addition to this, as models can be directly expressed in program code, we can encode other structures to make use of the full structure of a given language, which enables users to integrate PPSs into their standard model development process. One powerful *class* of models are simulators.

Simulators arise in many industrial [6, 15, 16] and scientific applications [24, 3, 10] encode years, if not decades, of research and development. Thus, Simulators by their very nature are information rich and naturally describe complicated joint densities. However, in order to write a simulator in a standard probabilistic programming set-up one would have to re-write a simulator in the language of the PPS, which in many cases is not feasible, nor efficient due to both the size and complexity of the simulator code base. However, what if one could arbitrarily connect a simulator, written in any language and connect it directly to a PPS, without having to re-write the simulator in the given PPS?

Our method of hijacking simulators builds on the work of [2] and extends the framework to encapsulate a more diverse range of simulators, in particular population-based simulators which arise in

a myriad of domains [BG: add citations]. Our framework provides a simple solution and makes it easy to transform arbitrary stochastic population-based simulators into probabilistic programs, regardless of the complexity of the simulator. Our system supports simulators written in 13 different programming languages and enables the user to perform for posterior inference over the density of the simulator. This enables two things. 1) It enables software developers and researchers to understand complicated code bases. 2) It provides interpretable inference results that provide policy makers with predictions that are truly interpretable, in the sense that the end-user of the inference results understands what physical events led to inference outcome. This is critical in several domains, and is particularly important in the medical domain. We provide examples of both in Section **TO ADD REF**.

In this paper, we introduce a novel method that allows one to shed light on the inner workings of a large class of population-based stochastic simulators. We achieve this by extending the work of [2] by interpreting such population-based simulators as probabilistic generative models within the framework of universal probabilistic programming (UPP) [20]. To this end, we *hijack* existing simulators by overriding their internal random number generators. Specifically, by replacing the existing low-level random number generator in a simulator with a call to a purpose-built UPP “controller”, which can thus control, track and manipulate the stochasticity of the simulator.

This allows for a variety of tasks to be performed on the hijacked simulator, such as running inference (by conditioning the values of certain draws and manipulating others), uncovering stochastic structure, and automatically producing result summaries, such as establishing the probability of different program paths/traces. By providing a common abstraction framework for different simulators, our approach further allows for easy and direct comparison between related or competing simulators, a characteristic that is valuable in the context of epidemiology simulators [8]. We provide a case study of the above in Section 4.2.

Our framework already supports application to simulators written in 13 industrial programming frameworks, and is easily extensible. This is crucial as, given the enormous code size and complexity, rewriting epidemiology simulators using a dedicated universal probabilistic programming language, such as Pyro [4], is often infeasible.

In time, we hope our approach will play a critical role in bringing recent advancements in probabilistic programming to bear on the vast array of existing simulators used throughout the sciences, thereby providing wide-ranging impacts across a number of fields.

This paper first gives an overview of existing Malaria simulators (Section ??), and proceed by introducing the necessary background on the *pyprob* framework and the concept of universal probabilistic programming (Section ??). Our approach is then demonstrated and analysed in the context of a Malaria case study (Section 4.2).

2 Background

- Paragraph on Probabilistic programming
- Paragraph on Simulators keep high-level focus on pop sims later
- Paragraph on Baydin et al

In-silico simulators have become a crucial tool in evidence-based decision-making within a large number of disciplines, including statistical physics [18], financial modeling [14], weather prediction [7], epidemiology [24] and many others. In many cases, simulation output can augment or even replace real data that may otherwise be costly or even impossible to generate. Recent advances in hardware have enabled simulations to model increasingly complex systems. Epidemiology studies the prevalence and spreading of diseases across populations. Recent advances in hardware have enabled simulations to model the dynamics of infectious diseases, such as Malaria, in ever greater detail.

Probabilistic programming [13, 26, 17] can be used to express probabilistic models and consequently perform automated inference these. Once a probabilistic model has been expressed in a probabilistic programming language, a wide range of inference techniques, such as Markov Chain Monte Carlo (MCMC)[9], Variational Inference (VI)[27] and Deep Neural Networks (DNN)[11], can be used by non-experts in an automated fashion.

3 Hijacking Simulators

3.1 The Engineering

Hijacking a simulator describes the process by which a simulator’s random number generators are replaced by calls to external sampling procedures, which are controlled by a probabilistic programming system (PPS). In practice, this amounts to performing a small number of surgical incisions into the simulator’s source code in order to replace built-in calls to random number generators. E.g., given a simulator written in C++/Boost[23], a Gaussian distribution object `boost::normal` is replaced with the corresponding `pyprob_cpp` distribution object, namely `pyprob_cpp::distributions::Normal`. Sampling from this distribution is then done by requesting the PPS to send a sample back to the simulator. **[BG: This will be extended]**

Algorithm 1 How to perform posterior inference over a stochastic simulator with

```

1: function CONVERT SIMULATOR TO A PROBABILISTIC PROGRAM( $A$ ) ▷
2:   replace raw samples statements  $f_t(x_t|\eta_t)$ 
3:   Create thin wrapper
4:   Create Containerized Environment  $C_i$ 
5:   Create Connection
6:   Run Program
7: end function
8: function CREATE THIN WRAPPER( $f_t(x_t|\eta_t)$ )
9:   to add
10: end function
11: function CREATE CONTAINERIZED ENVIRONMENT ( $\Delta_i$ ) ▷ Where  $\Delta_i$  is one of  $i$ -th
    dependencies
12:   to add
13: end function

```

The PPS and the simulator exchange xTensor objects <https://xtensor.readthedocs.io/en/latest/> through TCP or ICP using a generic FlatBuffers <http://google.github.io/flatbuffers/> protocol. On the PPS side, sampling is done using the deep learning framework PyTorch [22].

After a variable has been sampled, it is sent to the PPS using `pyprob_cpp::sample`, alongside with its simulator source trace. This allows the PPS to construct sample trace probabilities and other summary statistics (cf. Section 4.2).

Finally, the entry point to the simulator, i.e. `main` in C++, is replaced by a special `forward` call, in this case `pyprob_cpp::forward`. This allows the PPS to generate rollouts from the simulator remotely.

Population-based simulators In contrast to the setup used by [2], who interface a particle physics simulator that creates a single trace per `forward` call, a population-based simulator creates a trace per population member. This means that e.g. in EMOD, a standard scenario simulation rollout over the span of 3 years with a population of size $n = 5000$ [24] will generate about *four terabytes* of raw trace data. Unlike with [2], this amount of data cannot possibly be kept in RAM, which makes *a posteriori* trace analysis very inefficient. In order to deal with the shortcomings of `pyprob` in a population-based simulator context, we therefore extend the framework to be able to do trace analysis on the fly.

By extending `pyprob` to handle population-based simulators, the PPS can now track all stochastic random variables that are created within the simulator, which then allows us to generate trace plots and path probabilities associated to the execution paths of the program. The corresponding increase of simulator transparency helps reinstate much needed trust between policymakers and evidence-based methods.

3.2 The Theory

[BG: Tom add theory stuff here]

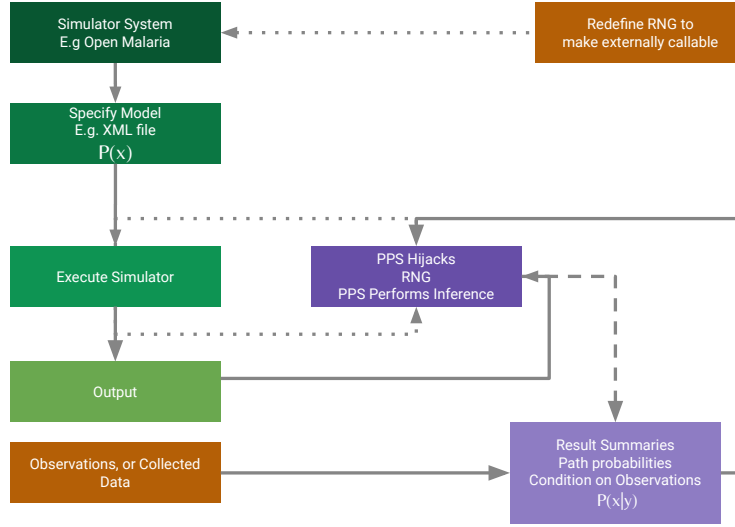


Figure 1: This flow-chart provides an overview of the process of how our system hijacks a generic population-based epidemiology simulator, such as EMOD, and how we modify the simulator to to hijack the random number generator (RNG).

4 Experiments

4.1 Epidemiology Simulators

Malaria epidemiology is governed by a complex set of drivers, few of which can be understood in isolation [5, 1, 24, 3]. These include within-host dynamics, population-specific traits and even local geography. Comprehensive modeling of all of these components remains challenging, particularly in a region-specific context. Computational epidemiology simulators have to reflect these complexities and are usually stochastic in nature. This can make simulation output highly non-trivial to interpret, particularly when trying to draw desired inferences coupled with observed data [21, 8].

We first demonstrate our system on a reduced version of the

Two the most advanced Malaria simulators, namely EMOD [3] and OpenMalaria [24], have proven to be particularly valuable to policymakers. OpenMalaria is based on microsimulations of Plasmodium falciparum in humans and was originally developed to simulate the impacts of malaria vaccines within simple villages or districts Compared with OpenMalaria, EMOD is able to simulate a variety of additional drivers, including complex geographies complete with migration and a large number of policy interventions Both EMOD and OpenMalaria are open source and implemented in C++.

[BG: Add linking stuff now to the toy-example where hopefully we can perform inference and]

4.2 Real-world Example

Ensemble methods are commonly used in statistics in order to combine the predictive power of multiple models [5, 25]. To this end, recent work has attempted to characterise the similarities and differences between two of the most advanced Malaria epidemiology simulators, EMOD [3] and OpenMalaria [24]. Evaluation is usually done by comparing a number of output parameters across a range of hand-crafted standard scenarios reflecting different geographical locations across Africa [25].

In the following, we illustrate how our method introduces a novel introspection paradigm. By extracting trace graphs from population-based simulators, policymakers can ask specific questions about properties of the model trace flow.

To illustrate the above, we present simulation output generated from a scenario resembling local conditions in the town of Ifakara (Tanzania). Both EMOD and OpenMalaria are configured to simulate a single population node of $n = 100$ and assume constant climatic conditions and no migration over the simulation period of 3 years. Please refer to Figure 2 for the seasonal Entomological Inoculation Rate (EIR), a measure of infectivity.




Figure 2: Seasonal Entomological Inoculation Rate (EIR) for the Ifakara scenario. Data is averaged over 30 day periods.

We provide examples of the generated trace plots from the connection between the simulator and the PPS in Figure 3.

We can see from the addressing schemes A_1, \dots, A_N (Tables 2 and 4.2) what physical events are connected to each other and how outputs in EMOD are generated from a different set of procedures as compared to OpenMalaria. By having access to such diagrams users can internally evaluate and scrutinize the decisions that the simulator is making.

This is important for policymakers, or general non-experts, as it not only details how we arrive at the given outputs, but it provides an understanding of which processes were most crucial in determining those outputs as can be seen from the path probabilities assigned to each of the vertices.

Additionally, by associating nodes in trace graphs representing the same physical processes within different models, the significance of model detail can be evaluated. For example, the full trace presented in Table 4.2 represent a sampling step associated with within-host dynamics of the Malaria parasite *Falciparum* in OpenMalaria node $A1$. The same physical process also occurs in EMOD's trace graph at position $A7$ (see Section 2). Comparisons like these could help developers better control model complexity, and even provide an alternative testing and debugging paradigm.

5 Discussions and Future Work

In this work we have demonstrated a method that enables one to hijack populationbased simulators, extending the work of [2]. We applied our method to two Malaria orientated population-based simulators and generated a variety of trace graphs. Finally, we have shown how our system enables policy makers and non-experts to analyse simulator outputs in a way previously unavailable in the field of epidemiology.

To extend our work further we aim to implement additional tools that will facilitate complicated inference procedures that condition on simulator output. We will also evaluate additional scenarios across Africa and South-East Asia to better understand the similarities and differences between EMOD and OpenMalaria.

References

- [1] Beatrice Autino, Alice Noris, Rosario Russo, and Francesco Castelli. Epidemiology of malaria in endemic areas. 4(1).

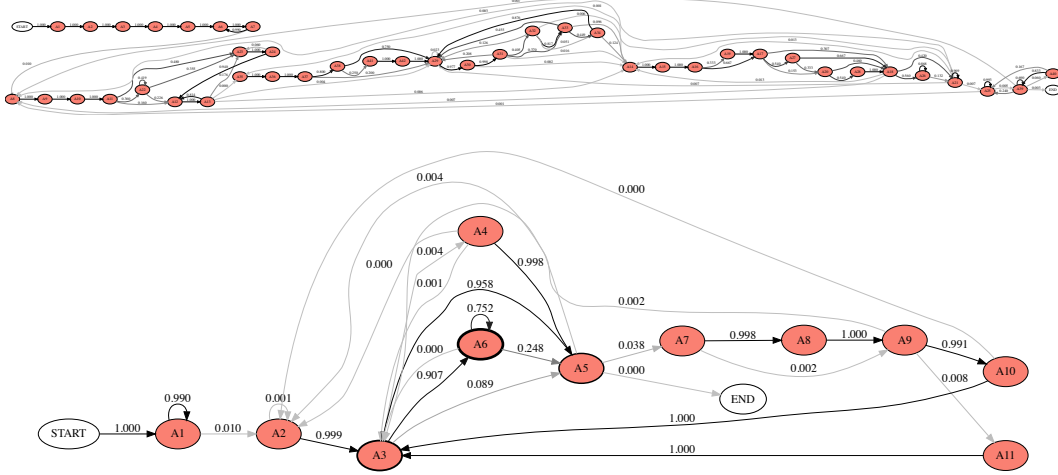


Figure 3: Here we run two equivalent models, compare the corresponding trace paths and corresponding path probabilities taken by the thousands of random variables generated internally within the simulators. **Top:** The specified model run in EMOD. **Bottom:** The specified model in OpenMalaria.

Table 1: An example of an address generated for the model run in the OpenMalaria simulator. We can see that A1 relates to Generating a member of the human population who may or may not be infect with the Malaria disease. We get something similar for EMOD, except this relates to A7 in the EMOD program execution.

Address ID	Full address
A1	[forward()+0x204;OM::Simulator::start(scenXml::Monitoringconst)+0x28a; OM::Population::createInitialHumans()+0x94; OM::Population::newHuman(OM::SimTime)+0x5c; OM::Host::Human::Human(OM::SimTime)+0x12b; OM::WithinHost::WHInterface::reateWithinHostModel(double)+0x99; OM::WithinHost::DescriptiveWithinHostModel::DescriptiveWithinHostModel(double)+0x3a; OM::WithinHost::WHFalciparum::WHFalciparum(double)+0xe6; OM::util::random::gauss(double, double)+0xb4]__Normal
⋮	⋮
A5	[forward()+0x204;OM::Simulator::start(scenXml::Monitoringconst&)+0x468; OM::Population::update1(OM::SimTime)+0xff; OM::Host::Human::update(bool)+0x2bc; OM::Clinical::ClinicalModel::update(OM::Host::Human&,double, bool)+0x96; OM::Host::NeonatalMortality::eventNeonatalMortality()+0x9; OM::util::random::uniform_01()+0xc0]__Uniform
⋮	⋮

- [2] Atilim Gunes Baydin, Lukas Heinrich, Wahid Bhimji, Bradley Gram-Hansen, Gilles Louppe, Lei Shao, Kyle Cranmer, Frank Wood, et al. Efficient probabilistic inference in the quest for physics beyond the standard model. *arXiv preprint arXiv:1807.07706*, 2018.
- [3] Anna Bershteyn, Jaline Gerardin, Daniel Bridenbecker, Christopher W Lorton, Jonathan Bloedow, Robert S Baker, Guillaume Chabot-Couture, Ye Chen, Thomas Fischle, Kurt Frey, et al. Implementation and applications of emod, an individual-based multi-disease modeling platform. *Pathogens and disease*, 76(5):fty059, 2018.
- [4] Eli Bingham, Jonathan P Chen, Martin Jankowiak, Fritz Obermeyer, Neeraj Pradhan, Theofanis Karaletsos, Rohit Singh, Paul Szerlip, Paul Horsfall, and Noah D Goodman. Pyro: Deep

Table 2: An interpretation table for each of the address of the overall trace generated from the corresponding OpenMalaria model.

Address ID	Interpretation
A1	Generate a human in the population within host dynamics
A2	Generate another human in the population within host dynamics
A3	The population is updated and a new human, or humans, may get infected
A4, A5	Potential child deaths within the population are simulated
A6	Determines parasite density of an individual infection
A7	Models how the disease is progressing within the infected humans
A8	Models how the disease is progressing within the population
A9	Models how the disease is progressing within the infected humans after the population has been updated
A10	Full clinical update on the population for those without severe or no Malaria infection.
A11	Full clinical update on the population for those with severe Malaria infections

universal probabilistic programming. *The Journal of Machine Learning Research*, 20(1):973–978, 2019.

- [5] Ewan Cameron, Katherine E Battle, Samir Bhatt, Daniel J Weiss, Donal Bisanzio, Bonnie Mappin, Ursula Dalrymple, Simon I Hay, David L Smith, Jamie T Griffin, et al. Defining the relationship between infection prevalence and clinical incidence of plasmodium falciparum malaria. *Nature communications*, 6:8170, 2015.
- [6] Valentina Di Pasquale, Salvatore Miranda, Raffaele Iannone, and Stefano Riemma. A simulator for human error probability analysis. *Reliability Engineering & System Safety*, 139:17–32, 2015.
- [7] Geir Evensen. Sequential data assimilation with a nonlinear quasi-geostrophic model using monte carlo methods to forecast error statistics. *Journal of Geophysical Research: Oceans*, 99(C5):10143–10162, 1994.
- [8] Casey Ferris, Benoit Raybaud, and Gregory Madey. OpenMalaria and EMOD: A case study on model alignment. In *Proceedings of the Conference on Summer Computer Simulation*, SummerSim ’15, pages 1–9. Society for Computer Simulation International. event-place: Chicago, Illinois.
- [9] Charles J Geyer. Practical markov chain monte carlo. *Statistical science*, pages 473–483, 1992.
- [10] T Gleisberg, Stefan Hoeche, Frank Krauss, Marek Schönherr, Steffen Schumann, Frank Siegert, and J Winter. Event generation with SHERPA 1.1. 2009.
- [11] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. The MIT Press, 2016.
- [12] Noah Goodman, Vikash Mansinghka, Daniel M Roy, Keith Bonawitz, and Joshua B Tenenbaum. Church: a language for generative models. *arXiv preprint arXiv:1206.3255*, 2012.
- [13] Andrew D Gordon, Thomas A Henzinger, Aditya V Nori, and Sriram K Rajamani. Probabilistic programming. In *Proceedings of the Future of Software Engineering*, pages 167–181. ACM, 2014.
- [14] P. Jäckel. *Monte Carlo Methods in Finance*. The Wiley Finance Series. Wiley, 2002.
- [15] Kenneth Judd, Lilia Maliar, and Serguei Maliar. Numerically stable stochastic simulation approaches for solving dynamic economic models. Technical report, National Bureau of Economic Research, 2009.
- [16] Michael Kass and Gavin Miller. Rapid, stable fluid dynamics for computer graphics. In *ACM Siggraph Computer Graphics*, volume 24, pages 49–57. ACM, 1990.
- [17] Dexter Kozen. Semantics of probabilistic programs. In *Foundations of Computer Science, 1979., 20th Annual Symposium on*, pages 101–114. IEEE, 1979.

- [18] David P. Landau and Kurt Binder. *A Guide to Monte Carlo Simulations in Statistical Physics*. Cambridge University Press, 4 edition, 2014.
- [19] Tuan Anh Le, Atilim Gunes Baydin, and Frank Wood. Inference compilation and universal probabilistic programming. *arXiv preprint arXiv:1610.09900*, 2016.
- [20] Tuan Anh Le, Atilim Güneş Baydin, and Frank Wood. Inference compilation and universal probabilistic programming. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 54 of *Proceedings of Machine Learning Research*, pages 1338–1348, Fort Lauderdale, FL, USA, 2017. PMLR.
- [21] Chikondi A. Mwendera, Christiaan de Jager, Herbert Longwe, Save Kumwenda, Charles Hongoro, Kamija Phiri, and Clifford M. Muteru. Challenges to the implementation of malaria policies in malawi. 19(1):194.
- [22] Adam Paszke, Sam Gross, Soumith Chintala, Gregory Chanan, Edward Yang, Zachary DeVito, Zeming Lin, Alban Desmaison, Luca Antiga, and Adam Lerer. Automatic differentiation in pytorch. In *NIPS-W*, 2017.
- [23] Boris Schäling. *The boost C++ libraries*. Boris Schäling, 2011.
- [24] T Smith, N Maire, A Ross, M Penny, N Chitnis, A Schapira, A Studer, B Genton, C Lengeler, Fabrizio Tediosi, et al. Towards a comprehensive simulation model of malaria epidemiology and control. *Parasitology*, 135(13):1507–1516, 2008.
- [25] Thomas Smith, Amanda Ross, Nicolas Maire, Nakul Chitnis, Alain Studer, Diggory Hardy, Alan Brooks, Melissa Penny, and Marcel Tanner. Ensemble modeling of the likely public health impact of a pre-erythrocytic malaria vaccine. 9(1):e1001157.
- [26] Sam Staton, Hongseok Yang, Frank Wood, Chris Heunen, and Ohad Kammar. Semantics for probabilistic programming: higher-order functions, continuous distributions, and soft constraints. In *Proceedings of the 31st Annual ACM/IEEE Symposium on Logic in Computer Science*, pages 525–534. ACM, 2016.
- [27] Martin J Wainwright, Michael I Jordan, et al. Graphical models, exponential families, and variational inference. *Foundations and Trends® in Machine Learning*, 1(1–2):1–305, 2008.
- [28] Frank Wood, Jan Willem Meent, and Vikash Mansinghka. A new approach to probabilistic programming inference. In *Artificial Intelligence and Statistics*, pages 1024–1032, 2014.
- [29] Yuan Zhou*, Bradley J Gram-Hansen*, Tobias Kohn, Tom Rainforth, Hongseok Yang, and Frank Wood. Lf-ppl: A low-level first order probabilistic programming language for non-differentiable models. *arXiv preprint arXiv:1903.02482*, 2019.